


TRIBHUVAN UNIVERSITY
INSTITUTE OF ENGINEERING
KATHMANDU ENGINEERING COLLEGE
Kalimati, Kathmandu



Project Report on
SPARK: Explainable Edge AI for Kinetic Pattern
Recognition and Distress Signaling

SUBMITTED BY:

Aaradhya Dev Tamrakar (79001)

Rupesh Kadel (79034)

Sankalpa Lamsal (79039)

Sonia Thapa (79043)

SUBMITTED TO:

DEPARTMENT OF ELECTRONICS, COMMUNICATION AND
INFORMATION ENGINEERING

July 2026

ACKNOWLEDGEMENT

We feel immense pleasure and are thankful for the opportunity to work on this project.

We would like to sincerely thank our supervisor, **Er. Dipen Manandhar**, for his unlimited encouragement and timely support and guidance in helping us with the different aspects of the project. We owe our profound gratitude to our project coordinator, **Er. Sujan Shrestha**, who took a keen interest in our project and guided us all along, providing the necessary information for developing a robust system.

We would like to thank **Er. Suramya Sharma Dahal**, Head of the Department of Electronics, Communication, and Information Engineering, for providing us with the opportunity to perform this project.

We are thankful to and fortunate enough to get constant encouragement, support, constructive criticism, and guidance from all teaching staff of the Department of Electronics, Communication, and Information Engineering.

ABSTRACT

SPARK is an open-hardware wearable fall detection system for Nepal’s growing elderly care gap. An ESP32 and MPU6050 IMU node, worn on the wrist or chest, runs a two-layer on-device pipeline a threshold pre-filter (Layer 1) gating a quantized INT8 TFLite Micro 1D CNN (Layer 2) with no cloud dependency and under 200 ms end-to-end latency. Confirmed falls trigger a Raspberry Pi gateway (FastAPI/PostgreSQL) that alerts caregivers via Telegram within 2 seconds, updates a Streamlit dashboard, and computes SHAP feature attribution per event so classifications are auditable rather than a black-box verdict, alongside auto-generated PDF incident reports.

The project contributes a Nepal-context wearable IMU fall dataset, a configurable sensitivity strategy (High / Standard / Sport), and an open-source reference design for low-cost elderly fall monitoring across South Asian care settings.

Keywords: Edge AI, elderly care, ESP32, fall detection, IMU, Nepal, SHAP explainability, TFLite Micro, wearable sensor.

TABLE OF CONTENTS

Acknowledgement	i
Abstract	ii
List of Abbreviations	v
List of Figures	vi
List of Tables	vii
List of Algorithms	1
1 INTRODUCTION	1
1.1 Background	1
1.2 Problem Statement	1
1.3 Objectives	2
1.4 Scope of the Project	2
1.5 Application of the Project	3
2 LITERATURE REVIEW	4
3 RELATED THEORY	6
3.1 Hardware Description	6
3.1.1 ESP32 DevKit V1	6
3.1.2 MPU6050 6-Axis IMU	6
3.1.3 Raspberry Pi 4B Gateway	7
3.1.4 Power System (Wearable Node)	7
3.1.5 Mechanical Mounting and Support Components	9
3.2 Software Description	11
3.2.1 Two-Layer Detection Architecture	11
3.2.2 1D CNN Architecture	11
3.2.3 SHAP Explainability Pipeline	11
3.2.4 Gateway Software Architecture	12
4 FEASIBILITY STUDY	13
4.1 Technical Feasibility	13
4.2 Economic Feasibility	14

5	METHODOLOGY	16
5.1	System Block Diagram	16
5.2	Hardware Implementation	18
5.2.1	Wearable Node Assembly	18
5.2.2	Power Budget Verification	18
5.3	Firmware Implementation (ESP32 PlatformIO)	19
5.3.1	IMU Acquisition and Filtering	19
5.3.2	Sliding Window Buffer	20
5.3.3	Layer 1 - Threshold Pre-filter	20
5.3.4	Layer 2 - TFLite Micro CNN Inference	20
5.4	Algorithm	20
5.4.1	WiFi-MQTT Event Transmission	22
5.4.2	Sensitivity Strategy Pattern	22
5.5	Machine Learning Pipeline	23
5.5.1	Dataset Preparation	23
5.5.2	1D CNN Training	23
5.5.3	INT8 Quantization and Firmware Integration	24
5.6	Gateway Software Implementation	24
5.6.1	Backend and Database	24
5.6.2	SHAP Attribution Pipeline	24
5.6.3	Telegram Alert	25
5.6.4	Streamlit Dashboard	25
5.6.5	PDF Incident Report	25
5.7	Testing and Evaluation Plan	25
5.7.1	Unit Testing	25
5.7.2	End-to-End Latency Measurement	25
5.7.3	ADL False-Positive Battery	26
5.7.4	CNN Performance Evaluation	26
6	EXPECTED OUTPUT	27
6.1	Deliverables	27
6.1.1	Hardware Deliverables	27
6.1.2	Software Deliverables	27
	References	30
	Appendix	32

LIST OF ABBREVIATIONS

ADL	Activities of Daily Living
API	Application Programming Interface
BEI	Bachelor of Electronics, Communication and Information Engineering
CNN	Convolutional Neural Network
ESP32	Espressif Systems ESP32 DevKit V1
IMU	Inertial Measurement Unit
INT8	8-bit Integer Quantization
IOE	Institute of Engineering
ISO	International Organization for Standardization
KEC	Kathmandu Engineering College
MQTT	Message Queuing Telemetry Transport
MPU6050	InvenSense MPU-6050 6-axis IMU
RPi	Raspberry Pi
SHAP	SHapley Additive exPlanations
SPARK	Signal Pattern Analysis & Real-time Kinetics
TFLite	TensorFlow Lite
WHO	World Health Organization

LIST OF FIGURES

3.1	ESP32 DevKit V1	6
3.2	MPU6050 6-Axis IMU	7
3.3	Raspberry Pi 4B Gateway	7
3.4	18650 Li-ion Cell	8
3.5	TP4056 Charger Module	8
3.6	AMS1117-3.3 LDO	8
3.7	INA219 Sensor	9
3.8	Velcro Wrist/Chest Strap	9
3.9	Chest Strap Mount	9
3.10	USB-C Cable	10
3.11	MicroSD Card (32 GB)	10
3.12	Resistor	10
3.13	Capacitor	11
5.1	System Block Diagram	16
5.2	Two-layer fall detection flowchart: Layer 1 threshold pre-filter (top zone) gates Layer 2 TFLite Micro CNN inference (bottom zone). Sensitivity profile thresholds are shown in the callout.	19
5.3	SPARK 1D CNN architecture: two Conv1D+MaxPool blocks compress the 200-sample IMU window to a 12-step feature map, which is flattened and classified by a Dense+Softmax head. INT8 quantized for TFLite Micro deployment.	24

LIST OF TABLES

4.1	Component Cost Estimation	15
6.1	Project Gantt Chart (Phase 3: September 2026 – March 2027)	29

Chapter 1: INTRODUCTION

1.1 Background

Falls are a leading cause of injury-related mortality and morbidity among adults aged 65 and above worldwide. The WHO Global Report on Falls Prevention in Older Age (2007) identifies falls as the second leading cause of unintentional injury death globally, with approximately 684,000 fatal falls occurring each year [1]. Non-fatal falls cause serious injuries including hip fractures, traumatic brain injuries, and long-term disability, placing a heavy burden on healthcare systems and caregivers.

Nepal is undergoing a rapid demographic shift. The Central Bureau of Statistics projects that Nepal’s population aged 65 and above will double by 2035 [2], driven by improving life expectancy and urbanization. Despite this growing demographic pressure, no domestic fall detection product exists in the Nepali market. Commercial solutions such as the Apple Watch and medical alert pendants are cost-prohibitive — a single Apple Watch SE (2nd generation) retails above NPR 30,000 — and require proprietary ecosystems unavailable to most Nepali families. Urban WiFi penetration in Kathmandu, however, is increasingly reliable, making a WiFi-based alert pipeline practically deployable in the target care settings.

Academic fall detection systems published in the literature typically rely on one of two approaches: threshold-only detection on raw accelerometer data, which produces unacceptably high false-positive rates on activities of daily living such as sit-to-stand transitions and stair climbing; or cloud-dependent machine learning inference, which introduces unacceptable network dependencies for a safety-critical application. Neither approach has been validated on a Nepal-context dataset, nor do any of the published open-source systems provide per-event classifier explainability for care staff.

SPARK addresses these gaps with a two-layer gated edge AI architecture running entirely on an ESP32 microcontroller, a SHAP-based explainability pipeline at the gateway, and a self-collected Nepal-context IMU dataset collected at KEC.

1.2 Problem Statement

Falls are the leading cause of injury-related mortality among adults over 65 globally [1]. Nepal’s elderly population is projected to double by 2035 [2], yet no affordable domestic fall detection solution exists. Commercial fall detection systems (Apple Watch, Samsung Galaxy Watch, Withings ScanWatch) are cost-prohibitive for Nepal’s eldercare context,

require proprietary ecosystems, and operate as closed black-box systems — returning a binary fall verdict with no classifier explanation. Published academic systems use either threshold-only detection (unacceptably high false-positive rate) or cloud-dependent machine learning inference (network dependency unacceptable for safety-critical applications in Nepal’s connectivity environment). No published open-hardware system demonstrates:

1. A two-layer gated architecture with on-device TFLite Micro inference.
2. SHAP-based per-event explainability surfaced to caregivers.
3. Adaptive sensitivity profiles.
4. Clinical-format automated incident reporting validated on a Nepal-context IMU dataset.

1.3 Objectives

This project focuses on building a wearable fall detection system that combines hardware, software, and edge AI to automate real-time monitoring and caregiver alerting. The scope includes the following key objectives:

1. To construct the hardware structure of a wearable edge computing node.
2. To integrate a two-layer gated AI architecture for cloud-independent, real-time fall classification.
3. To implement SHAP-based explainability techniques to log and render transparent feature attribution per event.

1.4 Scope of the Project

The scope of our project includes:

- **On-Device Data Acquisition:** We sample and filter 6-axis accelerometer and gyroscope data directly on the wearable node at a continuous 100 Hz frequency.
- **Dual-Layer Gated Inference:** Our system uses a fast mathematical threshold check to wake up an on-device, quantized 1D CNN classifier only when a potential fall signature is detected.
- **Gateway Ingestion & Storage:** We build a localized backend setup using FastAPI and PostgreSQL to securely save event history and user configuration data.

- **SHAP Explainability Pipeline:** The system calculates Shapley additive feature values at the gateway to show caregivers exactly which sensor movements triggered the alarm.
- **Instant Caregiver Alerting:** We route confirmed fall events to a Telegram Bot to message caregivers within two seconds of an incident.
- **Automated Clinical Reporting:** The pipeline automatically packages peak movement data and model confidence ratings into formal PDF incident summary files.

1.5 Application of the Project

1. Assistive Technology: Provides real-time, hands-free fall monitoring to increase the independence and safety of vulnerable individuals living alone.
2. Medical & Rehabilitation: Assists clinicians in tracking recovery progress and movement patterns during post-injury or post-stroke rehabilitation.
3. Elderly Care Facilities & Nursing Homes: Optimizes institutional care by allowing a small staff footprint to respond instantly to high-priority fall alerts.
4. Smart Home & IoT Integration: Functions as a localized, network-independent safety module that interfaces seamlessly with residential intranet hubs.
5. Clinical Documentation & Compliance Monitoring: Automatically logs and formats objective, auditable incident data required for formal medical reports and facility compliance audits.
6. Research & Development: Serves as an open-hardware reference design and provides a localized, South Asian regional IMU motion dataset for public academic research.

Chapter 2: LITERATURE REVIEW

Fall detection using wearable inertial sensors has been an active research area for over two decades. The literature can be broadly categorized by detection approach: threshold-based methods, machine learning classifiers, and more recent deep learning architectures.

Threshold-based methods compute a scalar magnitude from accelerometer axes — typically the resultant $|a| = \sqrt{a_x^2 + a_y^2 + a_z^2}$ and trigger a fall event when a spike followed by a quiet period is detected. Kangas et al. [3] compared multiple threshold algorithms on hip and wrist IMU data and reported sensitivity values of 73–95% depending on placement, but with false-positive rates exceeding 30% on vigorous ADLs. Bourke and Lyons [4] proposed a dual-axis threshold method achieving 96% sensitivity, though specificity on sit-to-stand transitions remained problematic. The fundamental limitation of threshold-only approaches is that any rapid motion — vigorous exercise, dropping an object, stumbling without falling — can trigger the threshold, requiring either a very high threshold (missing genuine falls) or a low threshold (excessive false positives).

Machine learning classifiers applied to hand-crafted time-domain and frequency-domain features from IMU windows have addressed some of these limitations. Sucerquia et al. published the SisFall dataset [5], a widely used benchmark containing 15 fall types and 19 ADL types from 38 subjects, enabling standardized evaluation of feature-based classifiers. Support vector machines, decision trees, and random forests trained on SisFall features have achieved sensitivity above 90%, but these approaches typically run on a connected laptop or server, not on the wearable hardware itself.

Deep learning on IMU time-series has demonstrated superior accuracy without manual feature engineering. Nho et al. [7] trained a 1D CNN on the UCI MHEALTH dataset and SisFall combined, reporting sensitivity of 97.2% and specificity of 98.1% with a model that fits within 150 KB. However, the model ran on a Raspberry Pi 3 serving as the wearable compute unit, not on a resource-constrained microcontroller. Wang et al. [8] applied an LSTM architecture to continuous IMU streams, achieving comparable accuracy but with inference latency unsuitable for real-time detection on embedded hardware.

TFLite Micro on microcontrollers has enabled CNN deployment on devices with fewer than 320 KB of RAM. David et al. [9] demonstrated keyword spotting and gesture recognition using TFLite Micro on Cortex-M4 class hardware, establishing the feasibility of INT8 quantized CNN inference in this memory envelope. Kumar et al. [10] deployed a threshold-gated classifier on an Arduino Nano 33 BLE Sense, but did not combine on-device TFLite

Micro CNN inference with a structured two-layer gating architecture or provide any gateway explainability.

Explainability in health AI has been identified as a critical gap for clinical adoption. Lundberg and Lee [11] introduced SHAP (SHapley Additive exPlanations), a unified framework for model interpretability based on cooperative game theory, which provides consistent attribution of each input feature’s contribution to a model output. While SHAP has been applied to ECG classification [12] and sepsis prediction in clinical ML, no published wearable fall detection system applies SHAP attribution at the per-event level and surfaces the result to caregivers in a clinical dashboard.

Gap summary: The existing literature leaves the following cluster of gaps unaddressed simultaneously: (a) on-device TFLite Micro CNN inference in a two-layer gated architecture on a sub-\$5 microcontroller; (b) SHAP-based per-event explainability surfaced to caregivers; (c) adaptive sensitivity profiles without firmware reflash; and (d) a dataset collected from Nepali subjects on locally procurable hardware. SPARK addresses all four gaps in a single open-source system.

Chapter 3: RELATED THEORY

3.1 Hardware Description

3.1.1 ESP32 DevKit V1

The ESP32 DevKit V1 (Espressif Systems) is a dual-core Xtensa LX6 microcontroller operating at up to 240 MHz, with 520 KB SRAM, 4 MB flash, and integrated 802.11 b/g/n WiFi and Bluetooth 4.2. For SPARK, only the WiFi subsystem is used; Bluetooth is disabled to conserve power. The ESP32 executes both the sensor acquisition ISR at 100 Hz and the two-layer fall detection pipeline including TFLite Micro CNN inference within the available SRAM and flash envelope. INT8 quantization reduces the 1D CNN model footprint to approximately 80–120 KB of flash, leaving sufficient headroom for the Arduino SDK, WiFi stack, and ring-buffer firmware.




Figure 3.1: ESP32 DevKit V1

3.1.2 MPU6050 6-Axis IMU

The InvenSense MPU-6050 integrates a 3-axis MEMS accelerometer ($\pm 2/\pm 4/\pm 8/\pm 16$ g selectable) and a 3-axis MEMS gyroscope ($\pm 250/\pm 500/\pm 1000/\pm 2000^\circ/\text{s}$ selectable) on a single die, communicating over I²C at up to 400 kHz. For SPARK, the accelerometer is configured at ± 8 g full-scale (captures the free-fall and impact spike of a genuine fall without clipping) and the gyroscope at $\pm 500^\circ/\text{s}$. The built-in Digital Motion Processor (DMP) is bypassed in favour of raw register reads at 100 Hz, enabling direct integration with the complementary filter and sliding window buffer.




Figure 3.2: MPU6050 6-Axis IMU

3.1.3 Raspberry Pi 4B Gateway

The Raspberry Pi 4B (4 GB) serves as the SPARK gateway. It runs the Mosquitto MQTT broker, the FastAPI/Uvicorn backend, PostgreSQL 14 event log, Streamlit dashboard, python-telegram-bot alert service, and the SHAP attribution pipeline. The 4 GB RAM variant is the target for this project; KEC has one such unit, with a dedicated purchase as contingency if it cannot be allocated. The RPi 4B is powered by a standard 5 V/3 A USB-C supply and can operate continuously from AC mains in a care facility or household setting. During development and demonstration, a laptop (minimum 4 GB RAM, dual-core CPU, Python 3.9+) can substitute as the gateway without code changes.




Figure 3.3: Raspberry Pi 4B Gateway

3.1.4 Power System (Wearable Node)

The wearable node is powered by a single 18650 lithium-ion cell (2000–3000 mAh typical) managed by a TP4056 charge controller module (USB-C input, 1 A charge current). An AMS1117-3.3 low-dropout linear regulator steps the cell voltage (3.0–4.2 V) to a stable 3.3 V rail for the ESP32 and MPU6050. An INA219 current/voltage sensor on the power rail provides real-time power telemetry logged to the gateway for battery life characterization. At approximately 80 mA active current draw, a 2000 mAh cell provides approximately 25 hours of continuous monitoring.

3.1.4.1 18650 Li-ion Cell

A single 3.7 V nominal protected lithium-ion battery (typically 2000–3000 mAh) that serves as the main power source for the wearable node. At an active current draw of approximately 80 mA, it provides around 25 hours of continuous, uninterrupted fall monitoring.



Figure 3.4: 18650 Li-ion Cell

3.1.4.2 TP4056 Charger Module

A single-cell constant-current/constant-voltage (CC/CV) linear charge controller module with USB-C input. It regulates a 1 A charging current to safely recharge the 18650 cell and provides safety cutoffs (overcharge protection at 4.2 V and over-discharge protection at 2.5 V).



Figure 3.5: TP4056 Charger Module

3.1.4.3 AMS1117-3.3

A low-dropout (LDO) linear voltage regulator that steps down the varying voltage of the 18650 cell (3.0 V to 4.2 V) to a fixed, stable 3.3 V power rail. This stable voltage is required to safely power the ESP32 microcontroller and the MPU6050 IMU sensor.



Figure 3.6: AMS1117-3.3 LDO

3.1.4.4 INA219 Sensor

A bi-directional current and voltage monitor module connected over I2C to the power rail. It provides real-time power telemetry, logging current draw and voltage levels to the gateway to help characterize battery life during continuous operation.




Figure 3.7: INA219 Sensor

3.1.5 Mechanical Mounting and Support Components

Beyond the core compute, sensing, and power modules, the wearable node relies on a set of mechanical mounting and low-level support components for physical attachment and circuit integration.

3.1.5.1 Velcro Wrist/Chest Strap

An adjustable Velcro strap used to secure the wearable enclosure to the wrist or chest of the user, allowing the accelerometer and gyroscope to reliably capture body motion during a fall event.



Figure 3.8: Velcro Wrist/Chest Strap

3.1.5.2 Chest Strap Mount

An elasticated chest-mount harness offering an alternative attachment point to the wrist strap, used during controlled fall data collection to evaluate sensor placement sensitivity on detection accuracy.



Figure 3.9: Chest Strap Mount

3.1.5.3 USB-C Cables

Standard USB-C cables used to power and program the ESP32 DevKit V1 boards and to charge the TP4056 module during bench testing.



Figure 3.10: USB-C Cable

3.1.5.4 MicroSD Card

A 32 GB MicroSD card used as the boot and storage medium for the Raspberry Pi 4B gateway, hosting the operating system, PostgreSQL event log, and locally cached SHAP attribution outputs.



Figure 3.11: MicroSD Card (32 GB)

3.1.5.5 Passive Components

Assorted resistors and capacitors used on the wearable node breadboard for signal conditioning, decoupling, and pull-up/pull-down configuration around the ESP32, INA219, and AMS1117 circuitry.



Figure 3.12: Resistor



Figure 3.13: Capacitor

3.2 Software Description

3.2.1 Two-Layer Detection Architecture

The detection pipeline implements a gated dual-layer structure. Layer 1 is a deterministic threshold pre-filter executed in the ESP32 ISR at 100 Hz. It computes the resultant acceleration magnitude $|a| = \sqrt{a_x^2 + a_y^2 + a_z^2}$ and the angular velocity magnitude $|\omega| = \sqrt{\omega_x^2 + \omega_y^2 + \omega_z^2}$. A PRE_FALL flag is raised when $|a|$ exceeds 2.5 g (free-fall proxy spike) followed by $|a| < 0.5$ g for more than 300 ms (post-impact rest phase). This threshold logic mirrors the deterministic safety logic of ISO 13482:2014 (safety requirements for personal care robots). Layer 1 latency is below 5 ms and incurs no inference cost.

When PRE_FALL is raised, Layer 2 is invoked: a 2-second sliding window of $(a_x, a_y, a_z, \omega_x, \omega_y, \omega_z)$ at 100 Hz is assembled (200 samples $\times$ 6 channels = 1200 INT8 values), normalized using StandardScaler parameters baked into the firmware at training time, and fed to the TFLite Micro 1D CNN. The CNN produces a fall probability score $P(\text{FALL}) \in [0, 1]$. If $P(\text{FALL}) > 0.85$, a CONFIRMED_FALL event is serialized as a JSON payload and transmitted to the gateway over WiFi via MQTT.

3.2.2 1D CNN Architecture

The 1D CNN consists of: Conv1D(32 filters, kernel 5) $\rightarrow$ ReLU $\rightarrow$ MaxPool1D(2) $\rightarrow$ Conv1D(64 filters, kernel 3) $\rightarrow$ ReLU $\rightarrow$ GlobalAveragePooling1D $\rightarrow$ Dense(32, ReLU) $\rightarrow$ Dense(2, Softmax). The model is trained in TensorFlow/Keras on UCI MHEALTH [6] and SisFall [5] datasets with self-collected Nepal-context data as a fine-tuning set. INT8 post-training quantization via the TFLite Converter produces a C array for inclusion in the PlatformIO build.

3.2.3 SHAP Explainability Pipeline

SHAP (SHapley Additive exPlanations) [11] assigns to each input feature i an attribution value ϕ_i satisfying local accuracy ($f(x) = \phi_0 + \sum_i \phi_i$), missingness, and consistency axioms. For each CONFIRMED_FALL event, the gateway computes SHAP values for the six IMU channel features (peak values of $a_x, a_y, a_z, \omega_x, \omega_y, \omega_z$ extracted from the 2-second

window) using the `shap` Python library’s `TreeExplainer` or `KernelExplainer` on the trained CNN output. The top contributing feature (e.g., `az_peak` with 42% of fall confidence) and the full SHAP vector are stored in the PostgreSQL `fall_events` table and displayed on the Streamlit dashboard alongside the fall event log. This provides care staff with an auditable rationale for every alert — a capability absent from all commercial and published academic fall detection systems.

3.2.4 Gateway Software Architecture

The gateway implements three software design patterns documented in the thesis Chapter 3 (OOSE alignment):

Factory Pattern decouples the detection node abstraction from its concrete implementations. A `FallDetector` abstract base class handles WiFi transmission, event serialization, and power management; `ThresholdDetector` (Layer 1) and `CNNDetector` (Layer 2) are concrete subclasses.

Observer Pattern implements the alert ecosystem. The FastAPI gateway is the Subject; Streamlit dashboard, Telegram bot, PostgreSQL logger, and PDF report generator are Observers subscribed to `CONFIRMED_FALL` events. Adding a new alert channel (e.g., SMS) requires adding a new Observer without modifying the gateway core.

Strategy Pattern implements the configurable sensitivity profiles. A `FallDetectionStrategy` interface is implemented by three concrete strategies: `HighSensitivity` (elderly, bed-ridden), `Standard` (default), and `SportMode` (active ADL). Threshold parameters are loaded from the user profile at gateway startup; no firmware reflash is required to switch profiles.

Chapter 4: FEASIBILITY STUDY

4.1 Technical Feasibility

SPARK is technically feasible because every component of the system has been individually demonstrated in the published literature, and the complete hardware stack is either in-hand or procurable locally in Kathmandu within one day.

On-device CNN inference on ESP32. TensorFlow Lite Micro has been demonstrated on Cortex-M4 class hardware with flash footprints below 250 KB [9]. The ESP32 Xtensa LX6 dual-core processor (240 MHz, 520 KB SRAM, 4 MB flash) provides a larger compute envelope than the Cortex-M4 targets evaluated by David et al. INT8 post-training quantization reduces the SPARK 1D CNN to approximately 80–120 KB of flash, leaving sufficient headroom for the Arduino SDK, WiFi stack, and ring-buffer firmware. The complementary filter, sliding window buffer, and Layer 1 threshold logic together add fewer than 4 KB of RAM and run at 100 Hz well within the ISR budget.

Two-layer gated detection architecture. Layer 1 threshold logic is a scalar comparison on the resultant acceleration magnitude $|a|$ and the post-impact rest duration Δt . Both computations complete in under 1 ms on the ESP32 at 100 Hz. Layer 2 CNN inference, when gated by Layer 1, runs only on candidate fall windows (typical rate: fewer than 5 per hour of normal ADL), keeping average power overhead negligible.

WiFi-MQTT pipeline. The ESP32 integrated 802.11 b/g/n transceiver with the PubSub-Client MQTT library has been demonstrated delivering JSON payloads to a Mosquitto broker on a local network within 50–100 ms under typical indoor conditions. The Raspberry Pi 4B gateway running Mosquitto, FastAPI/Uvicorn, PostgreSQL 14, and Streamlit targets the 4 GB RAM variant for headroom during simultaneous SHAP computation and PostgreSQL operation; a laptop (minimum 4 GB RAM, dual-core CPU, Python 3.9+) can substitute as the gateway during development without code changes.

SHAP gateway pipeline. The shap Python library’s KernelExplainer or GradientExplainer computes feature attributions for a 6-feature CNN output on the RPi 4B in under 200 ms per event — well below the 2-second Telegram alert deadline. Per-event SHAP computation is triggered only on confirmed fall events, not continuously, so gateway CPU load is minimal.

Self-collected dataset. Controlled fall simulation on a crash mat is a standard data-collection protocol used by Sucerquia et al. [5] (38 subjects) and Kumar et al. [10]. Four team members plus volunteers recording on-campus fall simulations at KEC can achieve the 500-fall-event

/ 2000-ADL-window target within four weeks. No ethical approval is required for healthy adult volunteers performing low-risk controlled falls in a supervised laboratory setting.

No unsolved technical dependencies. All software libraries (TFLite Micro, FastAPI, PostgreSQL, Streamlit, python-telegram-bot, ReportLab, SHAP) are open source and actively maintained. No custom PCB fabrication, no import of restricted components, and no cloud service subscriptions are required.

4.2 Economic Feasibility

SPARK requires no imported components; all hardware is procurable from Himalayan Solution, Breadfruit Electronics, or Daraz (Kathmandu). The primary gateway target is a Raspberry Pi 4B (4 GB); KEC has one such unit available, though allocation is not guaranteed. As a contingency, a dedicated RPi 4B (4 GB) purchase is included in the BOM. Table 4.1 summarises the estimated procurement cost for two wearable nodes (primary and spare) plus a third contingency ESP32 board and development consumables. Detailed component specifications are listed in Table 4.1.

Component	Qty	Unit (NPR)	Total (NPR)	Reference
ESP32 DevKit V1	3	1,200	3,600	Himalayan Solutions
MPU6050 IMU module	3	500	1,500	Himalayan Solutions
18650 Li-ion cell (protected)	3	250	750	SLT Pvt.Ltd
TP4056 charger module	3	90	270	SLT Pvt.Ltd
AMS1117-3.3 LDO	5	105	525	Giga Nepal
INA219 module	3	800	2,400	Daraz
Breadboard + jumper wires	2 sets	325	650	SLT Pvt.Ltd
Velcro wrist/chest strap	2	500	1,000	RR-papers
Enclosure (3D-printed PLA, KEC Makerspace)	500 g	2/g	1,000	KEC Makerspace
MicroSD 32 GB	1	1,000	1,000	SLT Pvt.Ltd
RPi 4B power supply	1	550	550	SLT Pvt.Ltd
Raspberry Pi 4B (4 GB)	1	18,699	18,699	RoboNepal
USB-C cables	3	267	801	Daraz
Resistors / capacitors assortment	1 lot	600	600	Himalayan Solutions
Total (without RPi 4B 4 GB)			14,646	
Total (with RPi 4B 4 GB)			33,345	

Table 4.1: Component Cost Estimation

Chapter 5: METHODOLOGY

5.1 System Block Diagram




Figure 5.1: System Block Diagram

SPARK consists of two physical subsystems: a battery-powered wearable node and a gateway unit, connected over a local WiFi network. Figure 5.1 illustrates the top-level system architecture.

The wearable node (ESP32 DevKit V1 + MPU6050) acquires raw 6-axis IMU data at 100 Hz, applies a complementary filter for noise rejection, and runs the two-layer fall detection pipeline entirely on-device. On a confirmed fall, a JSON event payload is published to the gateway over WiFi via MQTT. The Raspberry Pi 4B gateway ingests the event, fires a Telegram caregiver alert, updates the Streamlit dashboard, logs the event with SHAP attribution to PostgreSQL, and auto-generates a PDF clinical incident report. Figure 5.1 illustrates the end-to-end data and control flow across all pipeline stages.

Figure 5.1 is organized into four zones connected by two link types solid arrows for sequential on-device processing, and a dashed WiFi boundary where the wearable node hands off to the gateway. Each block is described below.

Sensor zone (wearable node, solid fill hatch)

- **MPU6050 IMU** : reads all six raw accelerometer and gyroscope axes over I²C at 100 Hz (Section 5.3.1).
- **Complementary filter** :fuses the raw axes ($\alpha = 0.98$) to reject high-frequency sensor noise before the signal reaches the detection pipeline.
- **Sliding window** : buffers the filtered stream in a 200-sample circular ring buffer (2 seconds at 100 Hz), continuously overwritten until a candidate fall window is latched.

Edge inference zone (wearable node, grid fill)

- **Two-layer detection** : Layer 1 threshold pre-filter gates Layer 2 TFLite Micro CNN inference, exactly as formalized in Algorithm 5.4.
- **MQTT publish** : on CONFIRMED_FALL, serializes the event as JSON and publishes it to the spark/events topic, crossing the WiFi boundary to the gateway.

Gateway zone (Raspberry Pi 4B, dot fill)

- **Mosquitto broker** :the local MQTT broker subscribed to spark/events, receiving the event published by the wearable node.
- **FastAPI/Uvicorn** :Incorporates the incoming event from the broker and orchestrates the remaining gateway pipeline stages.
- **PostgreSQL**: persists the event log, including the raw event fields and downstream SHAP output.
- **SHAP attribution** : computes per-feature attributions on the six peak IMU channels in under 200 ms, identifying which channel most contributed to the fall classification.

Output zone (parallel, dash-dot outline)

- **Telegram** : dispatches a caregiver alert within 2 seconds of event ingestion.
- **Streamlit Dashboard** : displays the event timestamp, severity score, and Layer 2 CNN confidence in real time.
- **PDF incident report** : auto-generates a one-page clinical report containing the SHAP attribution chart and a care-staff signature block.

5.2 Hardware Implementation

5.2.1 Wearable Node Assembly

The wearable node is assembled on a breadboard for the prototype phase, transitioning to a compact perfboard layout prior to the March 2027 demonstration. The assembly sequence is:

1. Connect the MPU6050 to the ESP32 via I²C: SDA to GPIO 21, SCL to GPIO 22, operating at 400 kHz.
2. Power the MPU6050 from the ESP32 3.3 V rail.
3. Connect the 18650 cell output (3.0–4.2 V) to the TP4056 module (battery management), then to the AMS1117-3.3 LDO input.
4. Connect the INA219 current sensor in series with the 3.3 V power rail to the ESP32 VIN; connect INA219 I²C lines to GPIO 21/22 (secondary address 0x41).
5. House the assembly in a 3D-printed PLA enclosure (KEC Makerspace, Anycubic filament) with a Velcro wrist strap and flush-mounted USB-C charging port.

5.2.2 Power Budget Verification

The INA219 provides real-time current and voltage telemetry logged to the gateway. Target metrics: active current draw ≤ 80 mA at 3.3 V, providing ≥ 25 hours of continuous monitoring from a 2000 mAh cell. This will be measured and documented in the thesis Chapter 5 (Implementation).

5.3 Firmware Implementation (ESP32 PlatformIO)




Figure 5.2: Two-layer fall detection flowchart: Layer 1 threshold pre-filter (top zone) gates Layer 2 TFLite Micro CNN inference (bottom zone). Sensitivity profile thresholds are shown in the callout.

5.3.1 IMU Acquisition and Filtering

The MPU6050 is configured via I²C at startup: accelerometer full-scale ± 8 g, gyroscope full-scale $\pm 500^\circ/\text{s}$, DLPF cutoff 20 Hz. A timer interrupt service routine (ISR) fires at 100 Hz and reads all six raw axes via direct register access (bypassing the DMP). A complementary filter ($\alpha = 0.98$) fuses accelerometer and gyroscope to produce filtered angular velocity estimates:

$$\hat{\omega}_k = 0.98 \cdot (\hat{\omega}_{k-1} + \omega_{\text{raw},k} \cdot \Delta t) + 0.02 \cdot \theta_{\text{accel},k} \quad (5.1)$$

Raw accelerometer values pass through without additional filtering (the DLPF suffices for noise rejection above 20 Hz).

5.3.2 Sliding Window Buffer

A circular ring buffer stores the most recent 200 samples (2 seconds at 100 Hz) of all six channels ($a_x, a_y, a_z, \omega_x, \omega_y, \omega_z$) in SRAM. When Layer 1 fires a PRE_FALL flag, the buffer is copied to a local array, normalized using StandardScaler parameters (mean and standard deviation per channel) baked into the firmware as constant arrays at training time, and passed to the TFLite Micro interpreter.

5.3.3 Layer 1 - Threshold Pre-filter

The resultant acceleration magnitude and angular velocity magnitude are computed every ISR tick:

$$|a| = \sqrt{a_x^2 + a_y^2 + a_z^2}, \quad |\omega| = \sqrt{\omega_x^2 + \omega_y^2 + \omega_z^2} \quad (5.2)$$

The PRE_FALL flag is set when: (1) $|a|$ exceeds 2.5 g for at least 1 sample (free-fall impact proxy), and (2) $|a|$ subsequently remains below 0.5 g for more than 300 ms (post-impact rest phase). The threshold values (2.5 g, 0.5 g, 300 ms) will be empirically calibrated against self-collected fall simulations and are adjustable via the Strategy Pattern sensitivity profiles (Section 5.4.2).

5.3.4 Layer 2 - TFLite Micro CNN Inference

5.4 Algorithm

The following steps describe the workflow of the SPARK fall detection system:

1. START
2. Initialize all hardware components:
 - MPU6050 IMU (6-axis accelerometer + gyroscope)
 - ESP32 Microcontroller (520 KB SRAM, 4 MB Flash)
 - WiFi module for MQTT communication
 - Raspberry Pi 4B Gateway
 - Mosquitto MQTT broker
3. Calibrate the sensor:

- Read raw accelerometer and gyroscope values at rest.
 - Compute sensor bias/offset for calibration.
 - Store calibration constants for runtime correction.
4. Continuous sensor sampling loop:
- Sample MPU6050 at 100 Hz over I²C.
 - Apply complementary filter ($\alpha = 0.98$) to fuse accelerometer and gyroscope data and reject noise.
 - Push filtered sample into a circular sliding window buffer (200 samples).
5. Fall detection loop:
- (a) Check if sliding window buffer is full (200 samples collected).
 - (b) Apply Layer 1 threshold check on windowed acceleration/gyroscope data.
 - (c) If threshold is not exceeded: discard window, return to Step 4.
 - (d) If threshold is exceeded: pass window to Layer 2 (1D CNN classifier).
 - (e) CNN outputs $P(\text{fall})$ and $P(\text{ADL})$ via softmax.
 - (f) If $P(\text{fall}) > 0.85$: mark event as FALL DETECTED.
 - (g) Else: mark as ADL (Activity of Daily Living) and return to Step 4.
 - (h) Log detection result with timestamp on-device.
6. Event publishing and transmission:
- Package fall event as JSON payload (timestamp, confidence score, sensor snapshot).
 - Publish payload via MQTT to topic `spark/events` over WiFi.
 - Mosquitto broker on Raspberry Pi receives the payload (~50–100 ms latency).
7. Gateway event processing:
- FastAPI/Uvicorn service ingests the incoming MQTT event.

- Store event record in PostgreSQL event log.
 - Run SHAP attribution on the 6 input features (<200 ms) to identify which features drove the CNN's decision.
8. Alert and visualization dispatch (parallel):
- Push event and SHAP explanation to Streamlit dashboard for live visualization.
 - Send Telegram alert to caregiver/care facility (<2 s).
9. Incident report generation:
- Compile PDF report containing: timestamp, severity, CNN confidence score, SHAP chart, and signature block.
 - Store report and link it to the corresponding event log entry.
10. Return to Step 4 and continue continuous monitoring.
11. END (on system shutdown or power-off)

5.4.1 WiFi-MQTT Event Transmission

On `CONFIRMED_FALL`, the ESP32 serializes the event as a JSON payload using `ArduinoJSON` and publishes it to the MQTT topic `spark/events` on the local Mosquitto broker. The JSON structure includes: `event_id` (UUID), `timestamp` (ISO 8601 UTC), `layer1_flag`, `layer2_confidence`, `ax_peak`, `ay_peak`, `az_peak`, and `severity_score` (1–5, derived from $|a|$ magnitude). Total transmission latency from `CONFIRMED_FALL` flag to MQTT publish is under 50 ms over a local WiFi network.

5.4.2 Sensitivity Strategy Pattern

Three sensitivity profiles are implemented as configurable parameter sets loaded from a JSON config file at gateway startup and communicated back to the ESP32 via MQTT on profile change:

- **HighSensitivity** (elderly, bed-ridden): Layer 1 threshold $|a|_{\text{spike}} = 2.0$ g, rest duration 200 ms; Layer 2 threshold $P(\text{FALL}) > 0.75$.
- **Standard** (default): $|a|_{\text{spike}} = 2.5$ g, rest duration 300 ms; Layer 2 threshold $P(\text{FALL}) > 0.85$.
- **SportMode** (active ADL): $|a|_{\text{spike}} = 3.0$ g, rest duration 400 ms; Layer 2 threshold

$$P(\text{FALL}) > 0.90.$$

Switching between profiles requires no firmware reflash — the gateway updates the config file and publishes the new parameters to the ESP32 via a dedicated MQTT topic.

5.5 Machine Learning Pipeline

5.5.1 Dataset Preparation

Training data is sourced from three datasets:

- **UCI MHEALTH Dataset** [6]: 23 physical activities from 10 subjects, 3 body-worn IMUs at 50 Hz. Resampled to 100 Hz via linear interpolation. Fall-relevant activity windows extracted.
- **SisFall Dataset** [5]: 15 fall types + 19 ADL types from 38 subjects at 200 Hz. Downsampled to 100 Hz. All fall types used as positive class; all ADL types used as negative class.
- **Self-collected KEC Dataset**: Controlled fall simulations (forward trip, lateral fall, backward fall, sitting-to-floor) and ADLs (walking, running, stair climbing, sit-to-stand, bending) recorded with the SPARK node at 100 Hz. Target: ≥ 500 fall windows + ≥ 2000 ADL windows. This dataset will be published to GitHub as “KEC Wearable Fall Dataset in Nepal Context.”

All datasets are windowed at 2 seconds with 50% overlap. Per-channel StandardScaler normalization is applied; scaler parameters (mean, std per channel) are extracted and baked into the firmware as constant arrays.

5.5.2 1D CNN Training

The 1D CNN is trained in TensorFlow/Keras on the combined UCI MHEALTH + SisFall dataset, then fine-tuned on the self-collected KEC dataset using transfer learning (freeze convolutional layers, retrain Dense layers). Training protocol: Adam optimizer ($\text{lr} = 1\text{e-}3$), batch size 64, 50 epochs with early stopping (patience = 10 on validation loss), 80/10/10 train/validation/test split stratified by subject. Class weighting is applied to address the natural imbalance between fall events and ADL windows.

Evaluation metrics: Sensitivity (Recall for FALL class), Specificity (Recall for NON_FALL class), F1-score, and AUC-ROC on the held-out test set. Target: Sensitivity $\geq 90\%$, Specificity $\geq 90\%$. Figure 5.3 illustrates the CNN layer sequence and the progressive compression of the time axis.




Figure 5.3: SPARK 1D CNN architecture: two Conv1D+MaxPool blocks compress the 200-sample IMU window to a 12-step feature map, which is flattened and classified by a Dense+Softmax head. INT8 quantized for TFLite Micro deployment.

5.5.3 INT8 Quantization and Firmware Integration

Post-training INT8 quantization is applied via the TFLite Converter with a representative calibration dataset (200 randomly sampled training windows). The quantized TFLite model is exported as a C byte array using the `xxd -i` command and included in the PlatformIO build as a header file. Model size target: ≤ 120 KB flash.

5.6 Gateway Software Implementation

5.6.1 Backend and Database

The FastAPI/Uvicorn backend exposes REST endpoints for event ingestion, profile management, and report generation. A Mosquitto MQTT consumer thread runs within the FastAPI process, subscribing to `spark/events` and persisting each `CONFIRMED_FALL` event to the PostgreSQL `fall_events` table (schema defined in Chapter 3, Table 3.1). The SHAP attribution pipeline (Section 5.6.2) is invoked as a post-insert task, updating `shap_top_feature` and `shap_json` columns before the Telegram alert is dispatched.

5.6.2 SHAP Attribution Pipeline

For each confirmed fall event, the gateway extracts the six peak IMU channel values (three accelerometer axes a_x, a_y, a_z and three gyroscope axes $\omega_x, \omega_y, \omega_z$, each taken at peak magnitude) from the 2-second event window and computes SHAP attributions using the `shap` library against the trained CNN (loaded as a TFLite interpreter on the gateway). The `KernelExplainer` or `GradientExplainer` is used depending on benchmark latency on the RPi 4B. The top contributing feature (e.g., $a_{z,peak}$, contribution 42%) is stored in `shap_top_feature` and the full 6-element SHAP vector is stored in `shap_json` (JSONB). Both fields are displayed on the Streamlit dashboard per event and included in the auto-generated PDF report.

5.6.3 Telegram Alert

The python-telegram-bot library dispatches a Telegram message to the configured caregiver group within 2 seconds of CONFIRMED_FALL event ingestion. The message includes: patient identifier, event timestamp (ISO 8601 local time + UTC offset), fall severity score (1–5), Layer 2 CNN confidence (as percentage), and the SHAP top-contributing IMU channel.

5.6.4 Streamlit Dashboard

The dashboard provides: (1) a live accelerometer magnitude plot (last 10 seconds, 100 Hz WebSocket stream via FastAPI), (2) a fall event log table (most recent 50 events with timestamp, severity, confidence, SHAP top feature, and false-positive dismiss button), (3) a sensitivity profile selector (HighSensitivity / Standard / SportMode), and (4) a “Generate Report” button triggering PDF incident report generation for the selected event.

5.6.5 PDF Incident Report

ReportLab generates a one-page PDF clinical incident report per fall event, containing: institution header (KEC), patient identifier, event timestamp, fall severity score, Layer 2 CNN confidence score, SHAP top-contributing feature, full SHAP vector bar chart, and a signature block for care staff acknowledgement. The report format is modelled on the clinical incident reporting structure recommended by WHO [1] for fall event documentation.

5.7 Testing and Evaluation Plan

5.7.1 Unit Testing

Each subsystem is validated independently before integration: (1) MPU6050 ISR correctness verified by comparing 100 Hz output against known waveforms from a signal generator; (2) Layer 1 threshold logic verified by injecting synthetic $|a|$ sequences with known fall/non-fall labels; (3) Layer 2 TFLite Micro inference verified by comparing ESP32 INT8 outputs against Python FP32 reference on 50 held-out windows; (4) MQTT payload verified by Wireshark capture; (5) Telegram alert latency measured from MQTT publish to message receipt timestamp.

5.7.2 End-to-End Latency Measurement

A GPIO pin on the ESP32 is toggled at Layer 1 PRE_FALL and at CONFIRMED_FALL MQTT publish, captured on an oscilloscope to measure on-device pipeline latency. Telegram message receipt timestamp is recorded manually. Target: on-device ≤ 200 ms, Telegram alert ≤ 2 seconds from fall event.

5.7.3 ADL False-Positive Battery

One hundred ADL sequences (walk, run, stair-climb, sit-to-stand, bend, drop object, vigorous arm swing) are performed wearing the node. The false-positive rate at Standard sensitivity profile is measured and documented. Target: false-positive rate $\leq 10\%$ on the ADL battery.

5.7.4 CNN Performance Evaluation

The held-out test set (10% of combined dataset, stratified by subject) is used to compute Sensitivity, Specificity, F1-score, AUC-ROC, and confusion matrix. Results are reported for both the FP32 Keras model and the INT8 TFLite Micro model to quantify quantization accuracy loss.

Chapter 6: EXPECTED OUTPUT

This chapter describes the deliverables, performance targets, and project schedule for SPARK. The system is developed in three sequential phases: proposal and planning (June–July 2026), skill development and dataset preparation (July–August 2026), and core build, integration, and evaluation (September 2026 – March 2027).

6.1 Deliverables

SPARK produces the following tangible deliverables by the March 2027 demonstration deadline:

6.1.1 Hardware Deliverables

1. **Functional wearable node** : ESP32 DevKit V1 with MPU6050 6-axis IMU, 18650 LiPo battery, TP4056 charger module, AMS1117-3.3 LDO regulator and INA219 power telemetry sensor, mounted in a wrist or chest enclosure 3D-printed in PLA at KEC Makerspace.
2. **Gateway unit** : Raspberry Pi 4B (4GB, KEC unit or separately procured) running the complete SPARK gateway stack (Mosquitto, FastAPI/Uvicorn, PostgreSQL, Streamlit, python-telegram-bot, ReportLab, SHAP pipeline).
3. **Power characterization report** : INA219 telemetry measurements documenting active current draw, idle current draw, and estimated battery life (target: ≥ 25 hours from 2000 mAh cell at Standard sensitivity profile).

6.1.2 Software Deliverables

1. **ESP32 firmware** (PlatformIO/Arduino C++) : MPU6050 ISR at 100 Hz, complementary filter, sliding window ring buffer, Layer 1 threshold pre-filter, Layer 2 TFLite Micro INT8 CNN inference, ArduinoJSON MQTT publisher, and MQTT-based sensitivity profile receiver.
2. **Trained and quantized 1D CNN model** — INT8 TFLite model file (≤ 120 KB) achieving Sensitivity $\geq 90\%$ and Specificity $\geq 90\%$ on the held-out test set, integrated as a C header array in the PlatformIO build.
3. **Gateway backend** (Python 3.11+): FastAPI/Uvicorn REST API with MQTT consumer, PostgreSQL `fall_events` schema with SHAP columns, SHAP attribution

pipeline (per-event, < 200 ms latency on RPi 4B), Telegram caregiver alert (< 2 s from fall confirmation), Streamlit clinical dashboard, and ReportLab PDF incident report generator.

4. **Self-collected KEC Wearable Fall Dataset:** Minimum 500 labelled fall event windows and 2000 ADL windows recorded at KEC campus using the SPARK node, with subject metadata, published to GitHub under an open licence as “KEC Wearable Fall Dataset in Nepal Context.”
5. **GitHub repository :** Complete open-source repository containing firmware, gateway software, trained model, dataset, KiCad schematic (breadboard-level), and a reproducible README enabling replication of the full system.

Gantt chart

Table 6.1 maps all major tasks to a 32-week schedule spanning September 2026 through March 2027 (Phase 3 only). Weeks 1–8 correspond to September–October 2026; weeks 9–16 to November–December 2026; weeks 17–24 to January–February 2027; weeks 25–32 to March 2027. Filled cells (●) indicate scheduled activity.

Table 6.1: Project Gantt Chart (Phase 3: September 2026 – March 2027)

Task	W1	W2	W3	W4	W5	W6	W7	W8	W9	W10	W11	W12	W13	W14	W15	W16
Self-collected dataset collection	●	●	●	●	●	●										
CNN training (UCI+SisFall baseline)	●	●	●	●												
CNN fine-tuning (KEC dataset)					●	●	●	●								
INT8 quantization + ESP32 flash							●	●								
ESP32 firmware (Layer 1 + Layer 2)	●	●	●	●	●	●	●	●								
Hardware node assembly + enclosure	●	●	●	●												
FastAPI + PostgreSQL backend	●	●	●	●												
SHAP attribution pipeline					●	●	●	●								
Telegram alert bot					●	●										
Streamlit dashboard						●	●	●	●							
End-to-end integration + dry run								●	●							
ADL false-positive battery test									●	●						
Sensitivity profile tuning									●	●						
PDF report auto-generation										●	●					
KEC dataset published to GitHub											●					
Thesis Chapters 1–4 draft											●	●	●			
Full system integration test												●	●			
F1/AUC evaluation on test set												●	●			
Thesis Chapters 5–6 draft													●	●		
Conference paper draft														●	●	
Final demo preparation															●	●
Thesis submission + defence																●

Note: W1–W4 = September 2026; W5–W8 = October 2026; W9–W12 = November–December 2026; W13–W16 = January–March 2027.

REFERENCES

- [1] World Health Organization, *WHO Global Report on Falls Prevention in Older Age*. Geneva, Switzerland: World Health Organization, 2007.
- [2] Central Bureau of Statistics, Nepal, *National Population and Housing Census 2021: Population Projection Report*. Kathmandu, Nepal: National Statistics Office, Government of Nepal, 2021.
- [3] M. Kangas, A. Konttila, P. Lindgren, I. Winblad, and T. Jämsä, “Comparison of low-complexity fall detection algorithms for body attached accelerometers,” *Gait & Posture*, vol. 28, no. 2, pp. 285–291, 2008.
- [4] A. K. Bourke, J. V. O’Brien, and G. M. Lyons, “Evaluation of a threshold-based tri-axial accelerometer fall detection algorithm,” *Gait & Posture*, vol. 26, no. 2, pp. 194–199, 2008.
- [5] A. Sucerquia, J. D. López, and J. F. Vargas-Bonilla, “SisFall: A fall and movement dataset,” *Sensors*, vol. 17, no. 1, p. 198, 2017.
- [6] O. Baños *et al.*, “mHealthDroid: a novel framework for agile development of mobile health applications,” in *Proc. 6th Int. Work-Conf. Ambient Assisted Living and Daily Activities (IWAAL)*, 2014, pp. 91–98.
- [7] Y.-H. Nho, J. G. Lim, and D.-S. Kwon, “Deep convolutional neural network for fall detection using inertial sensor data from wearable devices,” *IEEE Sensors J.*, vol. 21, no. 22, pp. 26305–26315, 2021.
- [8] Y. Wang, M. Chen, X. Wang, R. H. M. Chan, and W. J. Li, “LSTM-based fall detection system using wearable sensors,” *IEEE Sensors J.*, vol. 20, no. 11, pp. 6107–6116, 2020.
- [9] R. David *et al.*, “TensorFlow Lite Micro: Embedded machine learning for TinyML systems,” in *Proc. Machine Learning and Systems (MLSys)*, vol. 3, 2021, pp. 800–811.
- [10] A. Kumar, V. Sharma, and R. Singh, “Microcontroller-based fall detection for elderly monitoring using a threshold-gated classifier on Arduino Nano 33 BLE Sense,” *Int. J. Embedded Syst. Appl.*, vol. 12, no. 1, pp. 1–12, 2022.
- [11] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017,

pp. 4765–4774.

- [12] G. Yao *et al.*, “Interpretation of electrocardiogram heartbeat by CNN and GRU,” *Computational and Mathematical Methods in Medicine*, vol. 2021, 2021.

APPENDIX

A.1 Self-Collected Dataset Protocol

The SPARK self-collected dataset is recorded at KEC campus using the SPARK wearable node (ESP32 + MPU6050 at 100 Hz). All four team members plus recruited volunteers act as subjects. The protocol defines four fall types and five ADL types as listed below.

Fall types (positive class):

- Forward trip fall — simulated stumble onto a crash mat
- Lateral fall (left) — sideways fall onto crash mat
- Lateral fall (right) — sideways fall onto crash mat
- Sitting-to-floor fall — controlled seat loss from a chair

ADL types (negative class):

- Walking (normal pace, indoor)
- Running (moderate pace, corridor)
- Stair climbing (ascending and descending, KEC stairwell)
- Sit-to-stand transition (chair, rapid)
- Bending to pick up an object

Each fall is repeated a minimum of 5 times per subject. ADL sequences are recorded in 2-minute continuous blocks. The node is worn on the dominant wrist during all recordings. Raw CSV output is labelled using a post-hoc annotation script keyed to the Layer 1 flag timestamp.

A.2 ESP32 Hardware Connections

Wearable Node Pin Connections

Component	ESP32 Pin	Function
MPU6050 SDA	GPIO 21	I ² C Data
MPU6050 SCL	GPIO 22	I ² C Clock
MPU6050 VCC	3.3 V rail	Power
MPU6050 GND	GND	Ground
INA219 SDA	GPIO 21	I ² C Data (address 0x41)
INA219 SCL	GPIO 22	I ² C Clock
INA219 VIN+	After LDO output	Current sense (series)
TP4056 OUT+	LDO AMS1117 IN	Battery regulated output
AMS1117-3.3 OUT	ESP32 VIN	3.3 V regulated supply

A.3 PostgreSQL Schema

```
CREATE TABLE fall_events (  
    event_id          UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
    user_id           VARCHAR(64) NOT NULL,  
    timestamp         TIMESTAMPTZ NOT NULL,  
    layer1_flag        BOOLEAN NOT NULL,  
    layer2_confidence  FLOAT NOT NULL,  
    fall_confirmed     BOOLEAN NOT NULL,  
    ax_peak            FLOAT,  
    ay_peak            FLOAT,  
    az_peak            FLOAT,  
    severity_score     INT CHECK (severity_score BETWEEN 1 AND 5),  
    shap_top_feature   VARCHAR(32),  
    shap_json          JSONB,  
    alert_sent         BOOLEAN DEFAULT FALSE,  
    false_positive_flagged BOOLEAN DEFAULT FALSE  
);
```